

STORAGE SYSTEMS PROJECT

GROUP 19



***Stosys
Team**
Choosing us for
presentation

***Us**
VSFS
Partial working M5

FTL DESIGN

Writes append to respective logical block page map

Reads looks up respective logical block page map initially. If not found, then retrieves data from mapped data zone

Logical zone 0	
Logical zone 1	
Logical zone 2	
Logical zone 3	

GC iterates over each logical block and merge their Page map and existing data zone to new data zone.

Page map

Page map

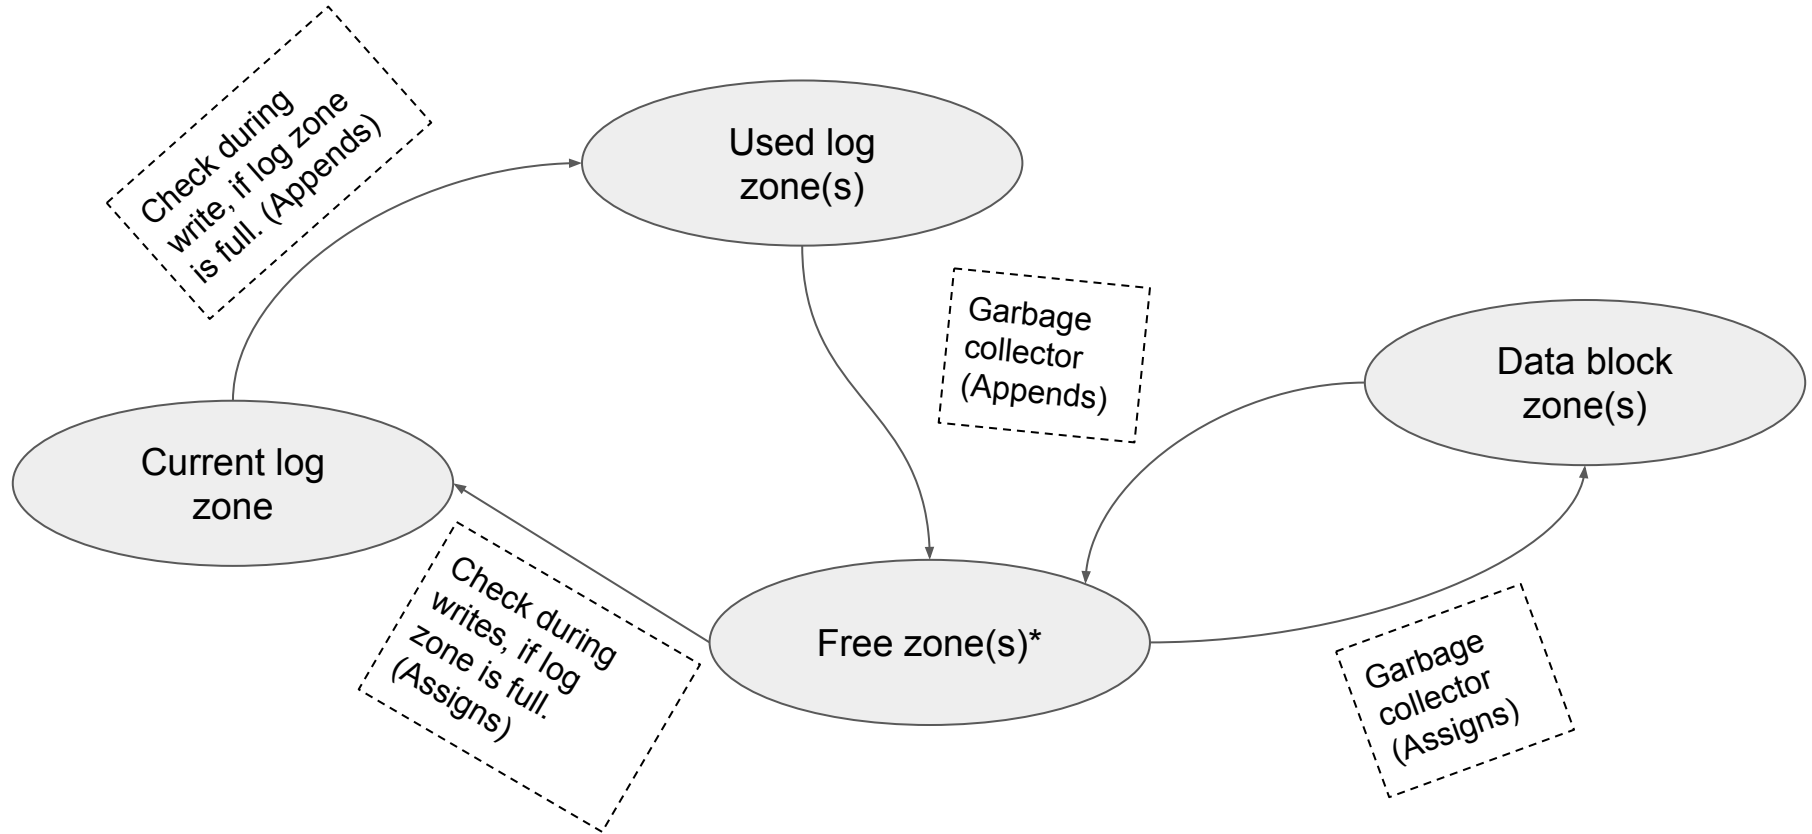
Page map

Data zone

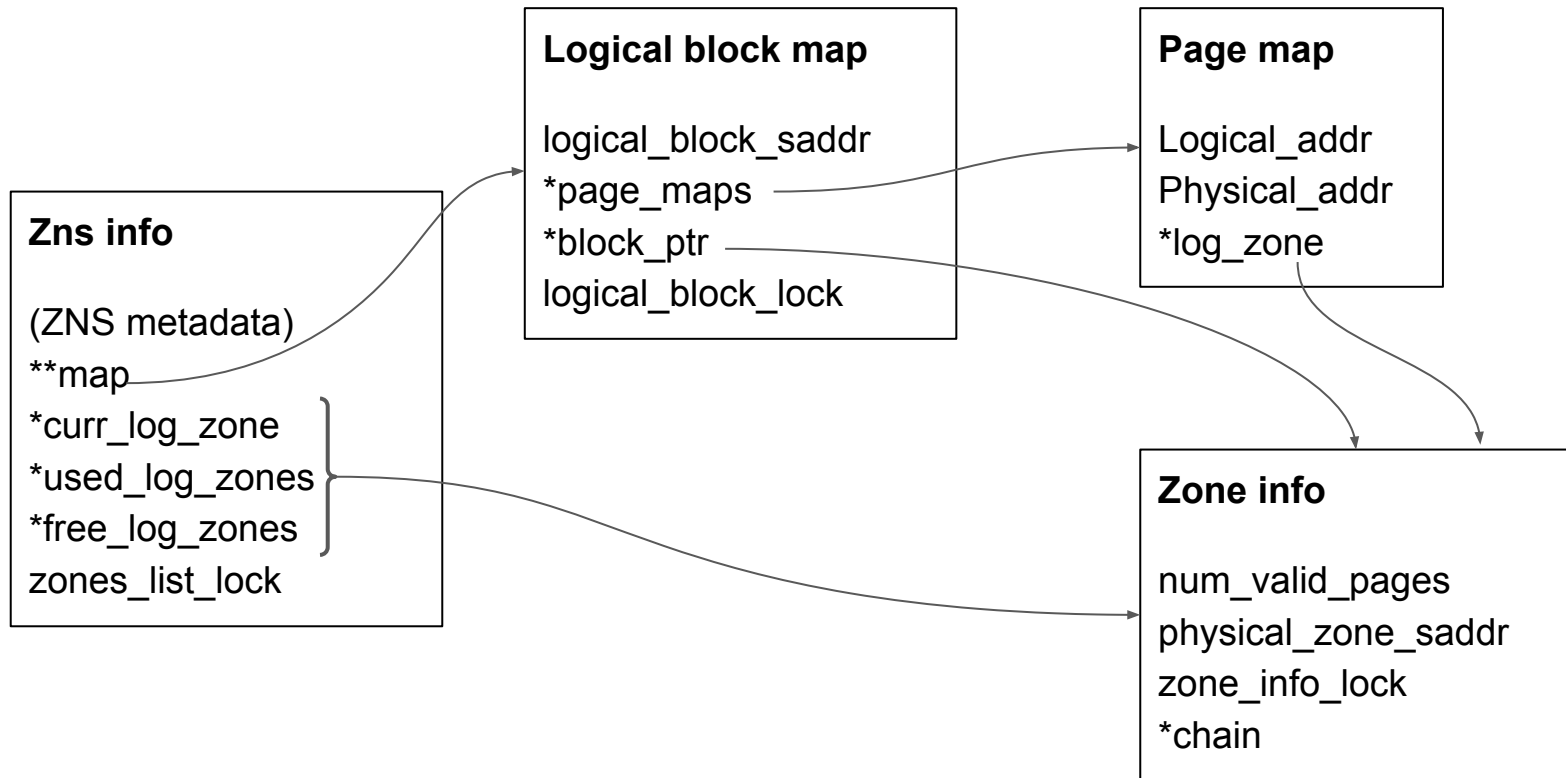
Data zone contains starting physical address of the zone

Page map contains logical address and respective physical address in log zone they are mapped to.

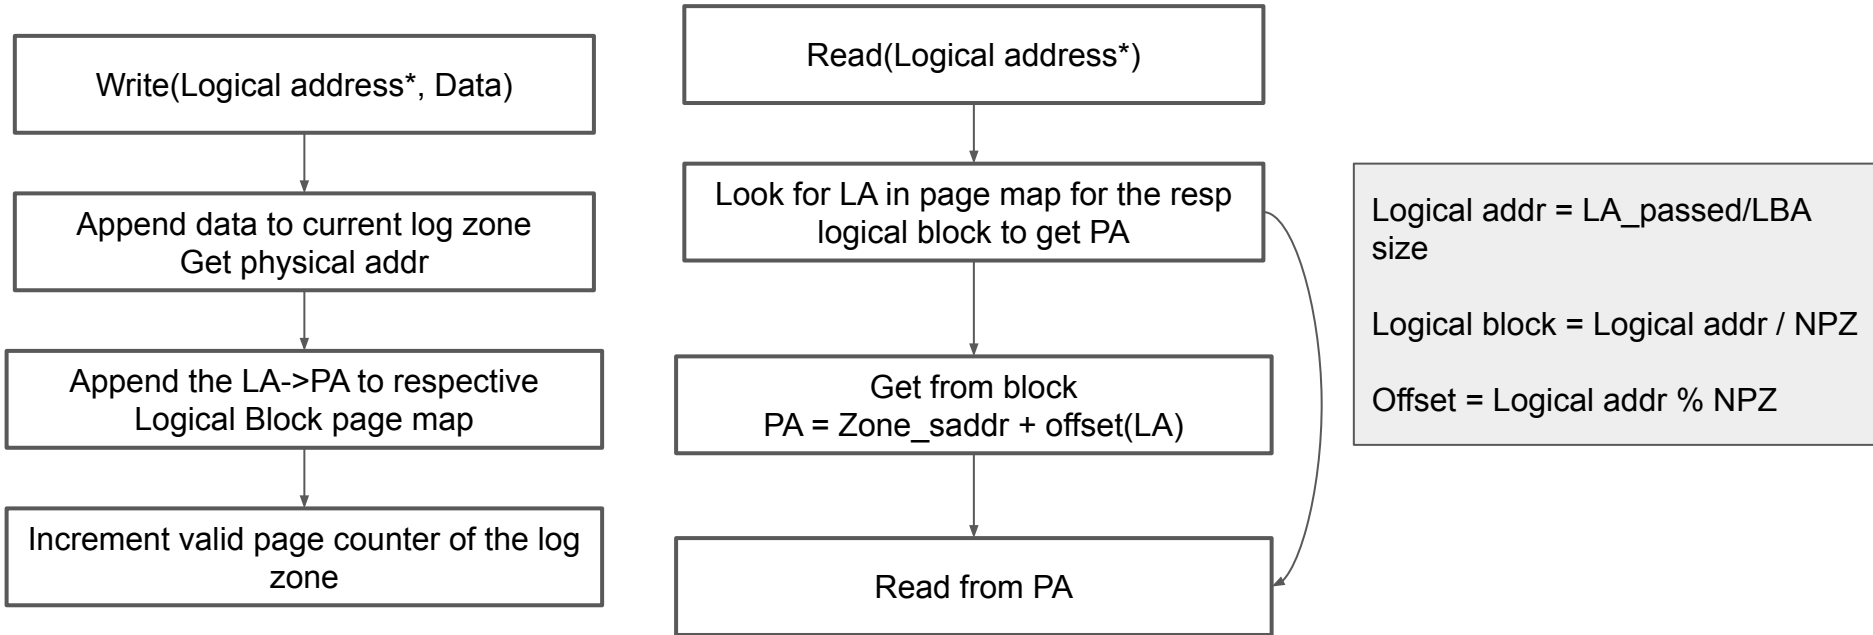
Zone State



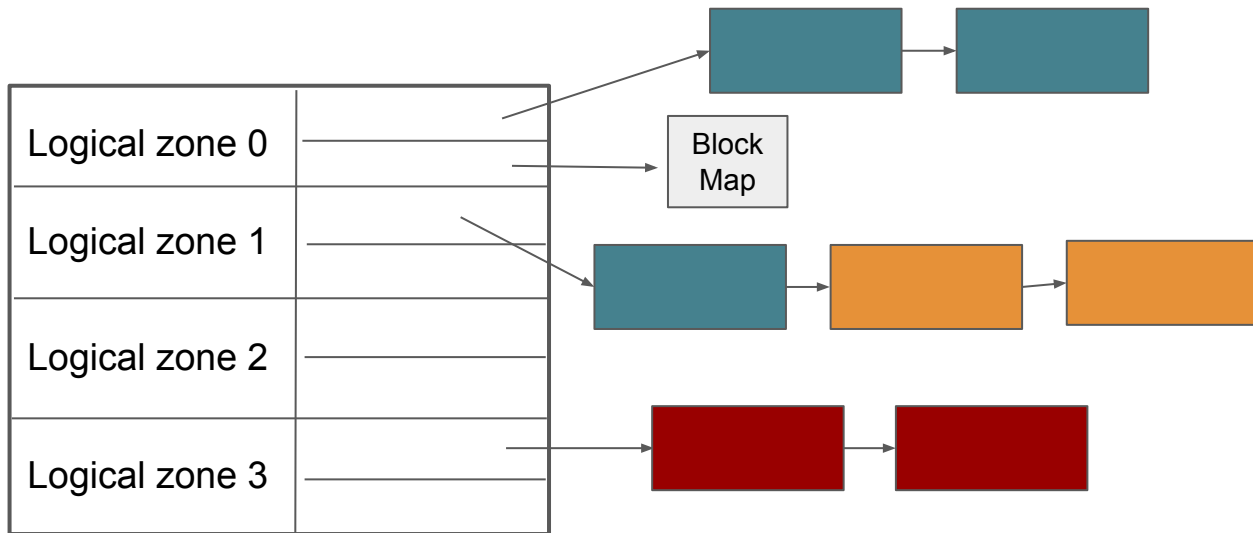
Data structures



Functionalities




Garbage collector working



Garbage Collector

```
For {  
  if(lzone->log != NULL)  
    lock(logical_zone#);  
    gc_log = lzone->log;  
    lzone->log = NULL;  
  unlock(logical_zone#);  
  nz = get_free_zone();  
  Merge(gc_log, lzone->block, nz);  
  Check_used_log_zones_to_free();  
  
  #Change lzone to next logical block  
}
```



FTL highlights

- Support to do any multiples of page size write
- Direct data zone append
- BitMaps to validate read is done after write

FILE SYSTEM DESIGN

Super block



□————□

8192
bytes

□————□

$254 * 4096$
bytes

*We followed very simple file system method with inplace updates.

Structures

```
Super Block {  
    bool Persistency;  
    uint32 InodePtr;  
    uint32 DataPtr;  
}
```

```
Inode {  
    uint32 Inode_no;  
    char entityName[235];  
    bool IsDir;  
    uint64 FileSize;  
    uint64 Indirect_lba;  
    uint64 Direct_lbas[480];  
}
```


```
Indirect_ptr {  
    uint64 addr;  
    uint64 Indirect_lba;  
    uint64 Direct_lbas[510];  
}
```

```
MYFS_Dir {  
    MYFS_DirData Entities[16];  
}
```

```
MYFS_DirData{  
    char EntityName[252];  
    uint32 inode;  
}
```

*Data structures are defined to be aligned with page size

In device



Main functions

Get_Path_Inode(FilePath, inode)

//If in mem

Look_Mem(FilePath, inode);

//If not

Get_Path_Inode(ParentPath, P_Inode)

Child_Inode# = Load_Children(P_Inode, Child_Name)

Addr = Child_Inode#*4096+SBlockOff

Read_From_NVM(Addr, inode);

Update Parent

CreateFile

CreateDirectory

RenameFile

DeleteFile

Persistence

FTL

Clear all data from log zone to data zone

Retrieve zone 0 to store block mapping

At start, check if mapping data is present in zone 0. If so, retrieve.

FILE SYSTEM

Flush all Inode data to NVM

Store Super Block with persistency byte set to true along with Bit maps

At start, Read Logical addr 0 retrieve super block and check if persistency set to true. If so, load bitmaps and root Inode.

If Time

- Make fully functional M5
- Log based File System
- More than one thread GC
- Observe how various parameters and intricate design changes affect performance.

Takeaways

- Never take Storage system and Software security together
- Storage system terminologies are brain teasers
- Extensive use of GDB for debugging
- Confidence to implement complex system
- Jump start to masters degree